

# COMP 5630/6630:Machine Learning

## Lecture 3: Linear Regression

- ML

- Learn a “mapping” from input to output  $h: X \rightarrow Y$ 
  - Fitting a function,  $f$ , from input to output using the *dataset*

- Supervised Learning


- Training Data
  - Used to learn a “mapping” from input to output,  $f: X \rightarrow Y$
- Hypothesis,  $h$ 
  - Learned “mapping”,  $h: X \rightarrow Y$  using the training data
- Test data
  - Use the hypothesis,  $h$  to predict new, unseen data

# House Price Prediction

- Given: a dataset that contains  $n$  samples

$$(x^{(1)}, y^{(1)}), \dots (x^{(n)}, y^{(n)})$$

- **Task:** if a residence has  $x$  square feet, predict its price?



15th sample  
 $(x^{(15)}, y^{(15)})$

$x = 800$


$y = ?$

# House Price Prediction

- Given: a dataset that contains  $n$  samples

$$(x^{(1)}, y^{(1)}), \dots (x^{(n)}, y^{(n)})$$

- **Task:** if a residence has  $x$  square feet, predict its price?



- Fit a line
- Estimate the price for  $x = 800$  on fitted line

*Is this the best way to learn a mapping from  $x \rightarrow y$ ?*


$$x = 800$$
$$y = ?$$

# House Price Prediction

- Given: a dataset that contains  $n$  samples

$$(x^{(1)}, y^{(1)}), \dots (x^{(n)}, y^{(n)})$$

- **Task:** if a residence has  $x$  square feet, predict its price?



$x = 800$   
 $y = ?$

- Linear vs quadratic fit

*How to choose between different mapping?*

# House Price Prediction

- Supervised Learning
  - Given a *training dataset*, learn a *mapping (hypothesis,  $h$ )* from  $x \rightarrow y$ , where  $y$  is labelled
  - Goal: Given a new datapoint,  $x$  (*test data*), predict the most accurate output,  $y$ , using the *learned hypothesis,  $h$* 
    - *learned mapping = trained model*



# Predicting House Price: Learn a Mapping from $x \rightarrow y$


Dataset of the living areas, bedrooms, and prices of 47 houses

Living area (ft <sup>2</sup> )	# bedrooms	Price (1000\$s)
1643	4	256
1356	3	202
1678	3	287
...	...	...
3000	4	400

# Predicting House Price: Learn a Mapping from $X \rightarrow y$

Dataset of the living areas, bedrooms, and prices of 47 houses

Living area (ft <sup>2</sup> )	# bedrooms	Price (1000\$s)
1643	4	256
1356	3	202
1678	3	287
...	...	...
3000	4	400



**X** **y**



# Recap: Linear Classifier


- Black and white circles are different labels.  $H_1, H_2, \dots$  represent different *decision boundaries* i.e. linear functions that best map the classification process.
  - **Goal:** find the best linear function that has highest accuracy

## Idea

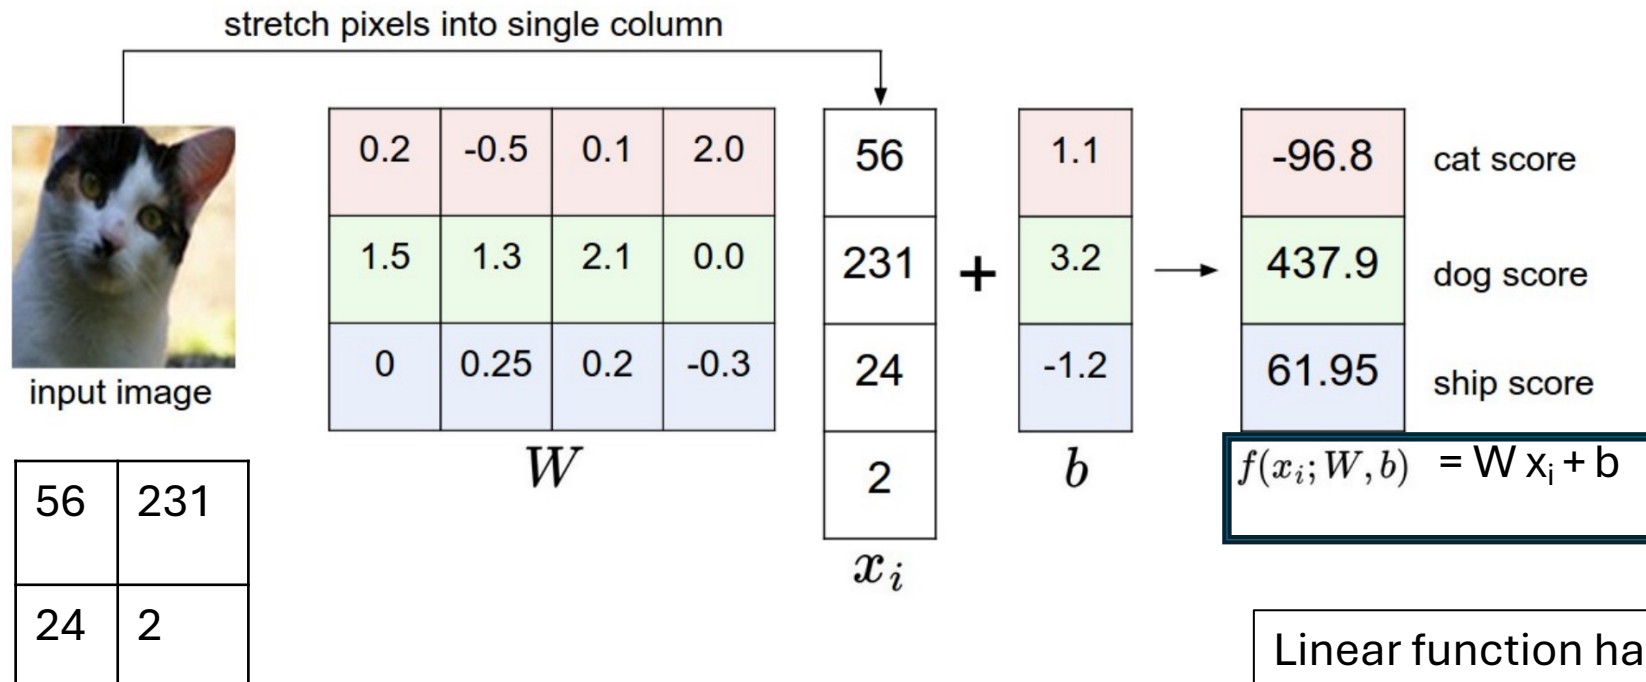
## Repeat

1. Draw a line to separate two classes
2. Calculate accuracy
3. Stop if accuracy can't be improved

Linear function to fit the data is a line  
 $y = mx + c$



# Recap: Linear Classifier Example



# Predicting House Price: Learn a Mapping from $X \rightarrow y$

Dataset of the living areas, bedrooms, and prices of 47 houses

Living area (ft <sup>2</sup> )	# bedrooms	Price (1000\$s)
1643	4	256
1356	3	202
1678	3	287
...	...	...
3000	4	400

$X$

$y$

- Predict  $y$  from one feature, living area

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1$$

$x_1$  = Living area  
 $x_2$  = # bedrooms

- Predict  $y$  from two features, living area and # bedrooms

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$

- $\theta_i$  = the **parameters** or **weights** of the linear model characterizing  $X \rightarrow Y$

- A more generic model is 
$$h(x) = \sum_{i=0}^n \theta_i x_i = \theta^T x,$$

# Linear Regression (cntd.)

- **Learning:** Given this formulation, we will need to identify a way to find the values of  $\theta$ .

$$h(x) = \sum_{i=0}^n \theta_i x_i = \theta^T x,$$

- We will need to use the ***training*** data to learn these parameters. This process is called ***learning***
- **What do we need to achieve this?**

# Linear Regression (cntd.)

- **Learning:** Given this formulation, we will need to identify a way to find the values of  $\theta$ .

$$h(x) = \sum_{i=0}^n \theta_i x_i = \theta^T x,$$

- We will need to use the **training** data to learn these parameters. This process is called **learning**
- **What do we need to achieve this?**
- We will define a function that measures the quality of predictions for each value of  $\theta$ .
- This is called the **cost function or objective function**

# How to Define the Cost or Objective Function

- Idea: Minimize the **squared** difference between the hypothesis,  $h_{\theta}(x)$  and  $y$

1.  $h_{\theta}(x) - y$  difference between the hypothesis,  $h_{\theta}(x)$  and  $y$
2.  $(h_{\theta}(x) - y)^2$  squared difference between  $h_{\theta}(x)$  and  $y$
3.  $\min_{\theta} (h_{\theta}(x) - y)^2$  choose  $\theta$  values to minimize the squared difference between  $h_{\theta}(x)$  and  $y$
4.  $\min_{\theta} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$  take the summation of squared difference of Step 3 for all training examples,  $m$
5.  $\min_{\theta} \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$  Multiply by a constant,  $\frac{1}{2}$ , for convention

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Objective function to minimize

**Ordinary least squares  
regression model**

# Learning: Gradient Descent

- Goal: To find a set of parameters  $\theta$  that will minimize the cost function  $J(\theta)$ .
- Common approach: *gradient descent*
- What does the *gradient descent* do?
  - Start with an initial “guess” for  $\theta$
  - Update values of  $\theta$  that will gradually move towards the “optimal solution”
  - What is the optimal solution?
  - The value of  $\theta$  that minimizes the cost function
- How do we do it computationally?

# Gradient Descent


- How do we do it computationally?

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

- $\alpha$  is the learning rate
  - Modulates how much of the change that we need to propagate at each instant
- Each update of  $\theta$  will be a step in the *steepest decrease of the cost function*  $J(\theta)$
- How to compute the derivative of the cost function  $J(\theta)$ ?




# Gradient Descent Visualization Example



Finding a (local) minimum  
depend on the starting point and  
step size,  $\alpha$ , to change direction

# Gradient Descent Visualization

In this example, depending on the starting point, the algorithm can end at P1 or P2. P2 is global minimum point.



# Gradient Descent Computation

$$\bullet \frac{\partial}{\partial \theta_j} J(\theta) = \frac{\partial}{\partial \theta_j} \frac{1}{2} (h_{\theta}(x) - y)^2$$

$$J(\Theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 \quad \text{Assume only 1 datapoint ignoring summation}$$

$$= 2 \cdot \frac{1}{2} (h_{\theta}(x) - y) \frac{\partial}{\partial \theta_j} (h_{\theta}(x) - y)$$

Chain rule of derivative

$$= (h_{\theta}(x) - y) \cdot \frac{\partial}{\partial \theta_j} (\theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n - y)$$

$$h(x) = \sum_{i=0}^n \theta_i x_i = \theta^T x,$$

$$= (h_{\theta}(x) - y) \cdot x_j$$

$$\frac{\partial}{\partial \theta_j} (\theta_j x_j) = x_j, \text{ other terms are 0}$$

e.g.,  $\frac{\partial}{\partial \theta_0} (\theta_0 x_0) = x_0, \text{ other terms are 0}$

# Gradient Descent

- Hence each update is given by

$$\theta_j := \theta_j + \alpha (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)}$$

- This is called the **LMS** update rule or the ***Least Mean Squares*** update rule.
  - Also known as the ***Widrow-Hoff*** learning rule.

# Gradient Descent

- Has several properties:
  - Magnitude of update is proportional to the error ( $y - h(x)$ )
    - What does this mean?
    - If we have a very good prediction i.e.  $h(x) \approx y$ , then the update is very small.
    - Conversely, if the prediction is very far off i.e.  $h(x) \gg y$  or  $h(x) \ll y$  then the update will be large.
- For learning over the complete training set, we iteratively update the parameters as below

Repeat until convergence {

$$\theta_j := \theta_j + \alpha \sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)} \quad (\text{for every } j).$$

}

# Gradient Descent Update: Batch vs Stochastic

- If we use all the examples in the training set **at once**:
  - Batch gradient descent
- What if we update at every single data point?
  - Stochastic or incremental gradient descent

```
Loop {  
    for i=1 to m, {  
         $\theta_j := \theta_j + \alpha (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)}$     (for every  $j$ ).  
    }  
}
```

# Summary


- If  $\alpha$  is too small: slow convergence.
- If  $\alpha$  is too large:  $J(\theta)$  may not decrease on every iteration; may not converge.

To choose, try  $\alpha$

$\dots, 0.001, \quad , 0.01, \quad , 0.1, \quad , 1, \dots$

# Normal Equations

Gradient Descent



Normal equation: Method to solve for  $\theta$  analytically.



# Normal Equations

- Steps:
  - Set the partial derivatives of the cost function  $J(\theta)$  to zero
  - Then we can estimate the parameters as follows:

$$\nabla_{\theta} J(\theta) = X^T X \theta = X^T \vec{y}$$

$$\Theta = (X^T X)^{-1} X^T y$$

- Where  $X$  is the input feature vector
- $y$  is the expected target value

# Normal Equation Example

$x_0$	$x_1$ = Living area (ft <sup>2</sup> )	$x_2$ = #bedrooms	y = Price(1000\$s)
1	1643	4	256
1	1356	3	202
1	1678	3	287

$$\bullet \mathbf{X} = \begin{bmatrix} 1 & 1643 & 4 \\ 1 & 1356 & 3 \\ 1 & 1678 & 3 \end{bmatrix}$$

$$\mathbf{y} = \begin{bmatrix} 256 \\ 202 \\ 287 \end{bmatrix}$$

Dimension of X,  $m \times n$  = m examples, n features  
Dimension of y =  $m \times 1$

$$\theta = (X^T X)^{-1} X^T y$$

**$m$  training examples,  $n$  features.**

## Gradient Descent

- Need to choose
- Needs many iterations  $\alpha$
- Works well even when  $n$  is large


## Normal Equation

- No need to choose
- Don't need to iterate  $\alpha$
- Need to compute  $(X^T X)^{-1}$
- Slow if  $n$  is very large.

# Basis Function: Extending Linear Regression to More Complex Models

- The inputs  $X$  for linear regression :
  - Original quantitative inputs
  - Transformation of quantitative inputs
    - e.g. log, exp, square root, square, etc.
  - Polynomial transformation
    - example:  $y = b_0 + b_1 \times x_1 + b_2 \times x_2 + b_3 \times x_3$

# Choice of Features



Blue curve fitting  $h_{\theta}(x) = \theta_0 + \theta_1(\text{size}) + \theta_2\sqrt{(\text{size})}$

Purple curve fitting  $h_{\theta}(x) = \theta_0 + \theta_1(\text{size}) + \theta_2(\text{size})^2$

# Basis Function: Extending Linear Regression to More Complex Models

- Previous
- $\theta_i$  = the **parameters** or **weights** of the linear model characterizing  $X \rightarrow Y$
- A more generic model is  $\phi$

$$h(x) = \sum_{i=0}^n \theta_i x_i = \theta^T x,$$

- Non-linear complex models

$$h(x) = \sum_{i=0}^n \theta_i \underbrace{\phi_i(x)}$$

Basis function

# Types of Basis Function


- Linear basis function: the simplest case

$$\varphi_i(\mathbf{x}) = x_i$$

- Polynomial basis function

$$\varphi_i(\mathbf{x}) = x^i$$

- A small change in  $x$  affects all basis functions




# Types of Basis Function

- Gaussian basis functions

$$\phi_j(x) = \exp \left\{ -\frac{(x - \mu_j)^2}{2s^2} \right\}$$

- A small change in x affects nearby basis functions.  $\mu_j$  and s controls location and slope






# Types of Basis Function

- Sigmoidal basis functions

$$\phi_j(x) = \sigma\left(\frac{x - \mu_j}{s}\right)$$

$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$

- A small change in x affects nearby basis functions.  $\mu_j$  and s controls location and slope



# Optional Reading

- Handout uploaded on Canvas