

COMP 5630/6630:Machine Learning

Lecture 4: Logistic Regression and Regularization

Predicting House Type: Learn a Mapping from $x \rightarrow y$

Dataset of the living areas, bedrooms, prices, and types of 47 houses

Living area (ft ²)	# bedrooms	Price (1000\$s)	Type
1643	4	256	Apartment
1356	3	202	Apartment
1678	3	287	House
...	
3000	4	400	House

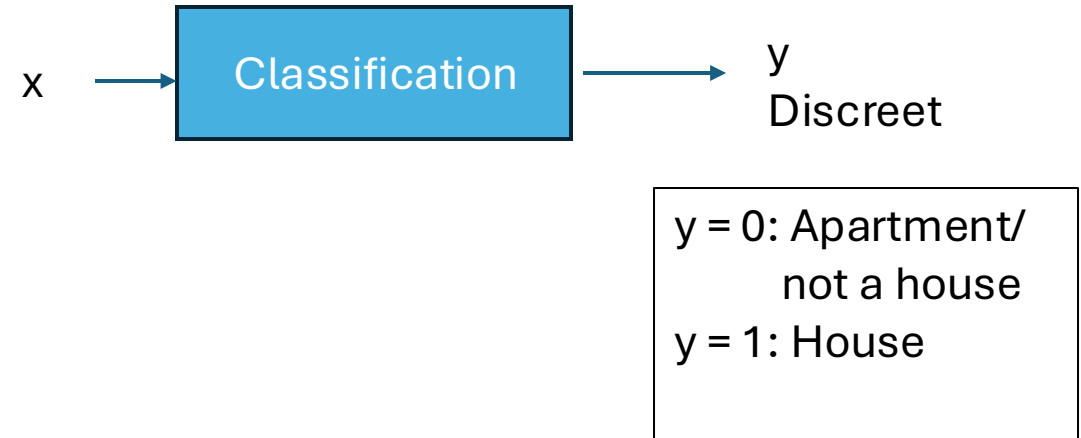
y = Class label
1 = House
0 = Apartment/ not a house

X = Input features

y = Class label

House Type Prediction

- Supervised Learning
 - Given a *training dataset*, learn a *mapping (hypothesis, h)* from $x \rightarrow y$, where y is labelled
 - Goal: Given a new datapoint, x (*test data*), predict the most accurate output, y , using the *learned hypothesis, h*
 - *learned mapping = trained model*



Logistic Regression: Hypothesis

- Recap: Linear regression hypothesis

$$h(x) = \sum_{i=0}^n \theta_i x_i = \theta^T x,$$

- *How does the logistic regression hypothesis differ from that of linear regression?*

Logistic Regression: Hypothesis

- Recap: Linear regression hypothesis

$$h(x) = \sum_{i=0}^n \theta_i x_i = \theta^T x,$$

- *How does the logistic regression hypothesis differ from that of linear regression?*
- Does not make sense to have predictions greater than 1 or less than 0.
 - Since the labels of y are 1 and 0 and $y \in \{0, 1\}$.

Logistic Regression: $0 \leq h_{\theta}(x) \leq 1$

Logistic Regression Model


- Goal
 - Define a function to satisfy $0 \leq h_{\theta}(x) \leq 1$
 - Take a probabilistic approach to represent $h = p(y=1 | \mathbf{x}; \boldsymbol{\theta})$
- Logistic regression model

$$h_{\theta}(\mathbf{x}) = g(\boldsymbol{\theta}^T \mathbf{x})$$

$$g(z) = \frac{1}{1 + e^{-z}}$$

$$h_{\theta}(\mathbf{x}) = \frac{1}{1 + e^{-\boldsymbol{\theta}^T \mathbf{x}}}$$


$$\boldsymbol{\theta}^T \mathbf{x} = \theta_0 + \sum_{j=1}^n \theta_j x_j$$



Logistic Regression Decision Boundary

$$h_{\theta}(x) = g(\theta^T x)$$


$$g(z) = \frac{1}{1+e^{-z}}$$



Suppose predict “ $y = 1$ ” if $h_{\theta}(x) \geq 0.5$


predict “ $y = 0$ ” if $h_{\theta}(x) < 0.5$

Logistic Regression Decision Boundary



Any Idea?

Logistic Regression Decision Boundary



Predict $y = 1$ if $-3 + x_1 + x_2 \geq 0$

Logistic Regression: Cost/ Objective Function to get Θ

- Gradient Descent!
- What do we need for gradient descent?
 - An objective function $J(\theta)$
 - A learning rate α
 - An initial “guess” for θ called θ_j
- Then, update θ_j until convergence as follows

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

$$\theta^T x = \theta_0 + \sum_{j=1}^n \theta_j x_j$$

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

Defining the Objective Function $J(\theta)$

- Let say, $P(y = 1 \mid x; \theta) = h_{\theta}(x)$ Chance of predicting a house given X
 $P(y = 0 \mid x; \theta) = 1 - h_{\theta}(x)$ Chance of predicting an apartment given X
Sum of two probabilities must equal to one

- Combining:

$$p(y \mid x; \theta) = (h_{\theta}(x))^y (1 - h_{\theta}(x))^{1-y}$$

Plugging $y = 1$ into $(h_{\theta}(x))^y (1 - h_{\theta}(x))^{1-y}$ yields $h_{\theta}(x)$

Plugging $y = 0$ into $(h_{\theta}(x))^y (1 - h_{\theta}(x))^{1-y}$ yields $1 - h_{\theta}(x)$

Defining the Objective Function $J(\theta)$

- Given: $p(y \mid x; \theta) = (h_\theta(x))^y (1 - h_\theta(x))^{1-y}$
- We want to estimate θ that will capture the dependency between y and x .
- We will instead call it the ***likelihood function*** that **maximizes** $p(y|\mathbf{x}; \theta)$ and is given by

$$L(\theta) = \prod_{i=1}^m p(y^{(i)} \mid x^{(i)}; \theta) = \prod_{i=1}^m (h_\theta(x^{(i)}))^{y^{(i)}} (1 - h_\theta(x^{(i)}))^{1-y^{(i)}}$$

Defining the Objective Function $J(\theta)$


- How can we get θ ?
- The principal of ***maximum likelihood***
 - Choose θ that makes the data as high probability as possible
 - Choose θ to maximize $L(\theta)$.
- Instead of maximizing $L(\theta)$, we can also maximize any strictly increasing function of $L(\theta)$.
 - We will maximize the log likelihood $\ell(\theta)$:

$$\ell(\theta) = \log L(\theta) = \sum_{i=1}^m y^{(i)} \log h(x^{(i)}) + (1 - y^{(i)}) \log(1 - h(x^{(i)}))$$

Recap: House Type Prediction Problem

- Given a *training dataset*, learn a *mapping (hypothesis, h)* from $x \rightarrow y$, where y is labelled

- Chance of predicting a house given X
- Chance of predicting an apartment given X



- $h \rightarrow$ *Choose θ* that will capture the dependency between y and x

$y = 0$: Apartment/
not a house
 $y = 1$: House

- Define likelihood
- Take log likelihood

- Goal of learning h :
 - Maximize log likelihood
 - Given a new datapoint, x (*test data*), predict the most accurate output, y , using the *learned hypothesis, h*

How do we get θ ?

- Gradient Descent!
- Then, update θ_j until convergence as follows

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

- So what is $\frac{\partial}{\partial \theta_j} J(\theta)$?

Gradient Descent $J(\theta)$

- Goal

- *Maximize* the log likelihood $= \sum_{i=1}^m y^{(i)} \log h(x^{(i)}) + (1 - y^{(i)}) \log(1 - h(x^{(i)}))$

- Equivalently, *minimize* $J(\theta)$


$$J(\theta) = \underbrace{-}_{\uparrow} \frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right]$$

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m \left[h_{\theta}(x^{(i)}) - y^{(i)} \right] x_j^{(i)}$$

Derivation is in the handout


Regularization

Example: Linear regression (housing prices)




$$\theta_0 + \theta_1 x$$

Underfit



$$\theta_0 + \theta_1 x + \theta_2 x^2$$

Just right




$$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$

Overfit


Overfitting: If we have too many features, the learned hypothesis may fit the training set very well, but fail to generalize to new examples (predict prices on new examples).

Example: Logistic regression




$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

(g = sigmoid function)



$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2 + \theta_5 x_1 x_2)$$




$$g(\theta_0 + \theta_1 x_1 + \theta_2 x_1^2 + \theta_3 x_1^2 x_2 + \theta_4 x_1^2 x_2^2 + \theta_5 x_1^2 x_2^3 + \theta_6 x_1^3 x_2 + \dots)$$

Addressing Overfitting


Options:

1. Reduce number of features.
 - Manually select which features to keep.
 - Model selection algorithm (later in course).
2. Regularization
 - Keep all the features, but reduce magnitude/values of parameters θ_j
 - Works well when we have a lot of features, each of which contributes a bit to predicting y


Intuition of Regularization



$$\theta_0 + \theta_1 x$$



$$\theta_0 + \theta_1 x + \theta_2 x^2$$



$$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$

Suppose we penalize and make θ_3 and θ_4 really small.

Minimize $J(\theta)$

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \boxed{1000 \theta_3 + 2000 \theta_4} \sim 0$$

Regularization in Linear Regression

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

$$\min_{\theta} J(\theta)$$

Normal Equation with Regularization

$$\bullet \mathbf{X} = \begin{bmatrix} 1 & 1643 & 4 \\ 1 & 1356 & 3 \\ 1 & 1678 & 3 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} 256 \\ 202 \\ 287 \end{bmatrix}$$

Dimension of X, $m \times n$ = m examples, n features

Dimension of y = $m \times 1$

x_0	x_1 = Living area (ft ²)	x_2 = #bedrooms	y = Price(1000\$)
1	1643	4	256
1	1356	3	202
1	1678	3	287

- Regularization: add an **extra Identity matrix** (n+1, n+1) to ($X^T X$) where n = # features, and extra 1 denotes the extra column of 1s for bias term.

$$\Theta = \left(X^T X + \lambda \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ \cdot & \cdot & \cdot & 1 & \cdot & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & 1 & \cdot & \cdot \\ \cdot & \cdot & \cdot & \cdot & \cdot & 1 & \cdot \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 \end{bmatrix}_{(n+1,n+1)} \right)^{-1} X^T Y$$

Non regularized: $\theta = (X^T X)^{-1} X^T y$

Logistic Regression with Regularization

- $\log L(\theta) + \lambda \sum_{j=1}^p |\beta_j|$ L1 penalty, known as “Lasso regression”
- $\log L(\theta) + \lambda \sum_{j=1}^p (\alpha \beta_j^2 + (1 - \alpha) |\beta_j|)$ L1 and L2 penalty, known as “Elastic net”