

COMP 5630/6630:Machine Learning

Lecture 8: MLP Demo, Project Guidelines

ML in Practice, Deep Learning

Today's Class

- ML in Practice
- Regularization in Neural Network in Practice
- Deep Learning

ML Pipeline

ML Pipeline

- Data Preprocessing
- Feature Engineering
- Model Selection
- Model Training
- Model Evaluation
 - Model Optimization

ML Pipeline

- Data Preprocessing
 - Feature Engineering
 - Model Selection
 - Model Training
 - Model Evaluation
 - Model Optimization
- Independent variables
 - Variable do not affected by other variables in the model
 - Dependent variable
 - Values are dependent on other variables
 - Example: prediction output

Data Preprocessing

- Data Cleaning
- Handling missing values and imbalanced data
- Feature Engineering
- Splitting
- Scaling and normalization

Data Preprocessing

- Data Cleaning
 - Check formatting
 - Example: strings to numeric categories, parsing dates

Data Preprocessing

- Handling missing values: identify missing values
 - `dataframe.isnull().values.any()`
 - returns True when there is at least one missing value occurring in the data
 - `df[df['height'].notnull]`
 - Returns whether the column has any rows with missing values

Data Preprocessing

- Handling missing values: handling missing values
 - Throw out the entire row if any value is missing
 - Replace missing values with mean/median

Data Preprocessing

- Handling imbalanced data
 - Most ML algorithms prefer balanced dataset. However, imbalanced data is common.
- Approaches to handle imbalanced data
 - Balanced random sampling
 - To make the data balance, randomly throw out instances from the majority class
 - Synthetic data generation
 - Generate synthetic data from the minority class. Example algorithm: SMOTE
 - Accept the imbalance dataset distribution 😊

Data Preprocessing

- Feature Engineering
 - Feature Selection
 - Relevant features for the task at hand
 - Reliable features to make prediction
 - Feature Construction
 - Refactor features
- Goal for feature selection
 - Eliminate spurious information
 - Enhancing efficiency
 - Reduce multicollinearity among features
 - multicollinearity : Feature variables are correlated in a model

Data Preprocessing

- Feature Selection
 - Statistical correlation
 - Identifying and removing features that are highly correlated with each other.
 - Information gain
 - Select features with high information gain
 - Automatic feature selection: scikit-learn library
 - **SelectKBest**: selects k best feature
 - **RFE (Recursive Feature Elimination)**: starts with all features, recursively eliminates features until a specified number of features remain.

Data Preprocessing

- Data splitting
 - Training/test/validation
 - K-fold cross validation

Data Preprocessing

- Data Scaling and Normalization
 - Z-score: **sklearn.preprocessing.StandardScaler**
 - Min max normalization: **sklearn.preprocessing.MinMaxScaler**
- Apply normalization on training set (not on test set)
- Apply learned normalization to test set

from sklearn.preprocessing

- *fit_transform()* on training data
- *transform()* on the test data

Model Selection

- Depends on the task at hand
 - Supervised learning: classification, regression
 - Example: predicting presence of a bug in a software
 - Predicting temperature
 - Unsupervised: Clustering
 - Example: Identifying patterns of behavior between high and low performing students
- Typically assess a series of models
 - For binary classification, one can evaluate logistic regression, decision trees, MLP, ...

Model Evaluation

- Regression

- Root Mean Squared Error (RMSE)
- Mean Absolute Error (MAE)
- The lower RMSE and MAE, the better is model

$$MAE = \frac{\sum_{i=1}^n |y_i - x_i|}{n}$$

$$RMSE = \sqrt{\sum_{i=1}^n \frac{(\hat{y}_i - y_i)^2}{n}}$$

- Classification

- Accuracy
- Precision
- Recall
- F1 score
- AUC-ROC
- The higher are the classification metrics, the better model performance.

- <https://www.aporia.com/learn/root-mean-square-error-rmse-the-cornerstone-for-evaluating-regression-models/>

Model Evaluation: Optimization

- Hyperparameter Tuning: Selecting the optimal set of parameters for ML model.
- Approaches for hyperparameter tuning
 - Grid Search: Systematically explores a predefined set of parameter values, effectively creating a grid of possible configurations.
 - Random search: similar to grid search, but instead of using all the points in the grid, it tests only a randomly selected subset of these points.
 - Example parameters
 - Neural network: Hidden layers, learning rate, kernels etc.
 - LR: solver, penalty, regularization

Model Evaluation & Optimization: Hyperparameter Tuning

- Hyperparameter Tuning Example: Logistic regression
- solver= {**'liblinear'**, **'newton-cg'**}
 - [depends on binary/multiclass LR. See documentation.]
- penalty = {L2, L1, None}
- Grid search= 2 x 3 = 6 parameter combinations
 - = {**'liblinear'**, **L1**} , {**'newton-cg'**, **None**} ,
 - = {**'liblinear'**, **L2**} , { **'newton-cg'**, **None**}
 - = {**'liblinear'**, **None**} , { **'newton-cg'**, **None**}

Model Evaluation & Optimization: Hyperparameter Tuning

- Hyperparameter Tuning Example: Logistic regression
- Grid search
 - Train LR with each of these 6 parameter combination.
 - Test on the validation set
 - Select the parameter combination which gives the best result
- Advantages
 - Finds the best combination of parameters
- Disadvantages
 - Exhaustive search, search space grows with number of parameters and values

Model Evaluation & Optimization: Hyperparameter Tuning

- Hyperparameter Tuning Example: Logistic regression
- Random search
 - Randomly select a subset (say, 4) parameter combinations from the 6 parameter combinations
 - Train LR with each of these parameter combination.
 - Test on the validation set
 - Select the parameter combination which gives the best result
- Advantages
 - Faster than grid search
 - The larger this dataset, the more accurate the optimization but the closer to a grid search.
- Disadvantages
 - The smaller this subset, the faster but less accurate the optimization.

Regularization in Practice: Neural Network


NN regularization?

Generalization error

- Performance on inputs not previously seen
 - Also called as *Test error*

Regularization is:

- “any modification to a learning algorithm to reduce its *generalization error* but not its *training error*”
- Reduce generalization error even at the expense of increasing training error
 - E.g., Limiting model capacity is a regularization method



Why regularization?

Generalization

- Prevent over-fitting

Occam's razor

Bayesian point of view

- Regularization corresponds to prior distributions on model parameters

Regularization: Limiting Number of Neurons

No. of input/output units determined by dimensions

Number of hidden units M is a free parameter

- Adjusted to get best predictive performance

Possible approach is to get maximum likelihood estimate of M for balance between under-fitting and over-fitting





Regularization: Limiting Number of Neurons

Sinusoidal Regression Prob

Two layer network trained on 10 data points

$M = 1, 3$ and 10 hidden units

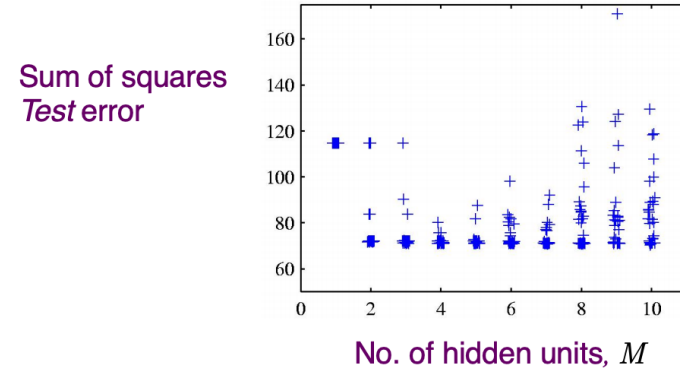
Minimizing sum-of-squared error function
Using conjugate gradient descent



Generalization error is not a simple function of M
due to presence of local minima in error function

Using Validation Set to fix no. of hidden units

Plot a graph choosing random starts and different numbers of hidden units M and choose the specific solution having smallest generalization error



- 30 random starts for each M
30 points in each column of graph
- Overall best *validation set* performance happened at $M=8$

There are other ways to control the complexity of a neural network in order to avoid over-fitting

Alternative approach is to choose a relatively large value of M and then control complexity by adding a regularization term


Parameter Norm Penalty

Generalization error not a simple function of M

- Due to presence of local minima
- Need to control capacity to avoid over-fitting
 - Alternatively choose large M and control complexity by addition of regularization term

Simplest regularizer is *weight decay*

$$\tilde{E}(\mathbf{w}) = E(\mathbf{w}) + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$



- Effective model complexity determined by choice of λ
- Regularizer is equivalent to a Gaussian prior over \mathbf{w}

In a 2-layer neural network

$$J_{\text{regularized}} = \underbrace{-\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(a^{[L](i)}) + (1 - y^{(i)}) \log(1 - a^{[L](i)}))}_{\text{cross-entropy cost}} + \underbrace{\frac{1}{m} \frac{\lambda}{2} \sum_l \sum_k \sum_j w_{kj}^{[l]2}}_{\text{L2 regularization cost}}$$

Regularization: Weight Decay

The name weight decay is due to the following

$$w_{t+1} = w_t - \alpha \nabla_w J - \lambda w_t$$

To prevent overfitting, every time we update a weight w with the gradient ∇J in respect to w , we also subtract from it λw .

This gives the weights a tendency to decay towards zero, hence the name.

Deep Neural Networks

What is Deep Learning?

1. Computational models composed of multiple processing layers
 - To learn representations of data with multiple levels of abstraction
2. Dramatically improved state-of-the-art in:
 - Speech recognition, Visual object recognition, Object detection
 - Other domains: Drug discovery, Genomics
3. Discovers intricate structure in large data sets
 - Using backpropagation to change parameters
 - Compute representation in each layer from previous layer
4. Deep convolutional nets: image proc, video, speech
5. Recurrent nets: sequential data, e,g., text, speech

Limitations of Conventional ML

- Limited in ability to process natural data in raw form
- Pattern Recognition and Machine Learning systems require careful engineering and domain expertise to transform raw data, e.g., pixel values, into a feature vector for a classifier

Automatic Representation Learning


- Methods that allow a machine to be fed with raw data to automatically discover representations needed for detection or classification
- Deep Learning methods are Representation Learning Methods
- Use multiple levels of representation
 - Composing simple but non-linear modules that transform representation at one level (starting with raw input) into a representation at a higher slightly more abstract level
 - Complex functions can be learned
 - Higher layers of representation amplify aspects of input important for discrimination and suppress irrelevant variations

Example: Images

- Input is an array of pixel values
 - First stage is presence or absence of edges at particular locations and orientations of image
 - Second layer detects motifs by spotting particular arrangements of edges, regardless of small variations in edge positions
 - Third layer assembles motifs into larger combinations that corresponds to parts of familiar objects
 - Subsequent layers would detect objects as combinations of these parts
- Key aspect of deep learning:
 - These layers of features are not designed by human engineers
 - They are learned from data using a general purpose learning procedure

Deep versus Shallow Classifiers

- Linear classifiers can only carve the input space into very simple regions
- Image and speech recognition require input-output function to be insensitive to irrelevant variations of the input,
 - e.g., position, orientation and illumination of an object
 - Variations in pitch or accent of speech
 - While being sensitive to minute variations, e.g., white wolf and breed of wolf-like white dog called Samoyed
 - At pixel level two Samoyeds in different positions may be very different, whereas a Samoyed and a wolf in the same position and background may be very similar



Selectivity-Invariance dilemma

- Shallow classifiers need a good feature extractor
- One that produces representations that are:
 - selective to aspects of image important for discrimination
 - but invariant to irrelevant aspects such as pose of the animal
- Generic features (e.g., Gaussian kernel) do not generalize well far from training examples
- Hand-designing good feature extractors requires **engineering skill and domain expertise**
- Deep learning learns features automatically

- <https://medium.com/@Tms43/understanding-padding-strides-in-convolutional-neural-networks-cnn-for-effective-image-feature-1b0756a52918>